



Bahir Dar University

Faculty of Computing
Department of Software Engineering
OSSP Individual Assignment

System Call Implementation

Full Name: Amanuel Fikadie Debeb

Student ID: BDU1600914

Section: A

System Call: nice()

Submitted To: Lecturer Wondimu B.

Submission Date:

System Call Implementation Example

nice() in User Space (glibc)

When you call `nice(int inc)` in C, you're using the glibc wrapper for the system call. This wrapper simplifies the interaction with the kernel. Here's how it works:

```
#include <unistd.h>
#include <sys/resource.h>
#include <errno.h>
int nice(int inc) {
    int ret;
    ret = getpriority(PRIO_PROCESS, 0) + inc;
    if (ret < -20) ret = -20;
    if (ret > 19) ret = 19;
    if (setpriority(PRIO_PROCESS, 0, ret) == -1) {
        return -1;
    }
    return ret;
}
```

Explanation:

The `nice()` function adjusts the priority of the calling process.

It retrieves the current nice value using `getpriority()`, adjusts it by `inc`, and enforces bounds (-20 to 19).

The new nice value is set using `setpriority()`.

Kernel Space: The Actual System Call

In the Linux kernel source code, the `nice()` system call is implemented in `kernel/sched/core.c`. Here's a simplified breakdown:

```
SYSCALL_DEFINE1(nice, int, inc) {
    long ret;
    ret = task_nice(current, inc);
    return ret;
}
```

Key Steps:

SYSCALL_DEFINE1: Macro to define a system call with one argument (`inc`).

task_nice(): Kernel function that adjusts the nice value of the current process (`current`).

Bounds Check: Ensures the nice value stays within the valid range (-20 to 19).

Under the Hood: task_nice()

The task_nice() function delegates the priority adjustment to the scheduler. Here's a simplified view:

```
long task_nice(struct task_struct *p, int inc) {
    int new_nice = PRIO_TO_NICE(p->static_prio) + inc;
    if (new_nice < -20) new_nice = -20;
    if (new_nice > 19) new_nice = 19;
    set_user_nice(p, new_nice);
    return new_nice;}
```

Explanation:

PRIO_TO_NICE: Converts the process's static priority to a nice value.

set_user_nice(): Updates the process's priority in the scheduler's data structures.

Summary of Layers

Layer	Function	Role
User Space	nice(int inc)	Glibc wrapper that adjusts process priority and handles bounds.
Syscall	SYSCALL_DEFINE1(nice)	Kernel entry point for the nice system call.
Scheduler	task_nice()	Adjusts the process's priority and interacts with the scheduler.

Practical Example

Writing a C Program to Use nice()

Here's a complete example demonstrating the nice() system call:

```
#include <stdio.h>
#include <unistd.h>
#include <sys/resource.h>
#include <errno.h>

int main() {
    int current_nice = getpriority(PRIO_PROCESS, 0);
    printf("Current nice value: %d\n", current_nice);

    int result = nice(10); // Lower priority
    if (result == -1 && errno != 0) {
        perror("nice failed");
        return 1;
    }

    int new_nice = getpriority(PRIO_PROCESS, 0);
    printf("New nice value: %d\n", new_nice);
}
```

```
// Simulate work
for (int i = 0; i < 5; i++) {
    printf("Working... %d\n", i + 1);
    sleep(1);
}
```

```
return 0;}
```

Steps to Compile and Run:

CTRL + O then ENTER and CTRL + X

Save the code to nice_example.c.

Compile: gcc nice_example.c -o nice_example.

Run: ./nice_example.

Expected Output:

Current nice value: 0

New nice value: 10

Working... 1

Working... 2

...

Key Takeaways

The nice() system call modifies process priority, influencing CPU scheduling.

Lower nice values give higher priority (more CPU time), while higher values reduce priority.

The kernel enforces bounds (-20 to 19) to prevent misuse.

Understanding nice() is essential for system programming and resource management.
